

Best practices for Cloud DLP

It can be difficult to figure out where Cloud DLP fits in your architecture or to identify requirements for Cloud DLP. Here are some best practices for you to understand how to use Cloud DLP in various scenarios:

- **Use data profiling versus inspection jobs:** Data profiling allows you to scan BigQuery tables in a scalable and automated manner without the need for orchestrating jobs. Considering the growth of data and the increasing number of tables, leveraging profiling features is recommended as it takes care of orchestration and running inspection jobs behind the scenes without any overhead. The inspection jobs can complement profilers when deeper investigation scans are needed. For example, if there are around 25,000 tables to be scanned, the recommendation is to scan all the tables with a profiler and then do a deep scan of 500 tables to flag sensitive/unstructured data that needs a more exhaustive investigation.

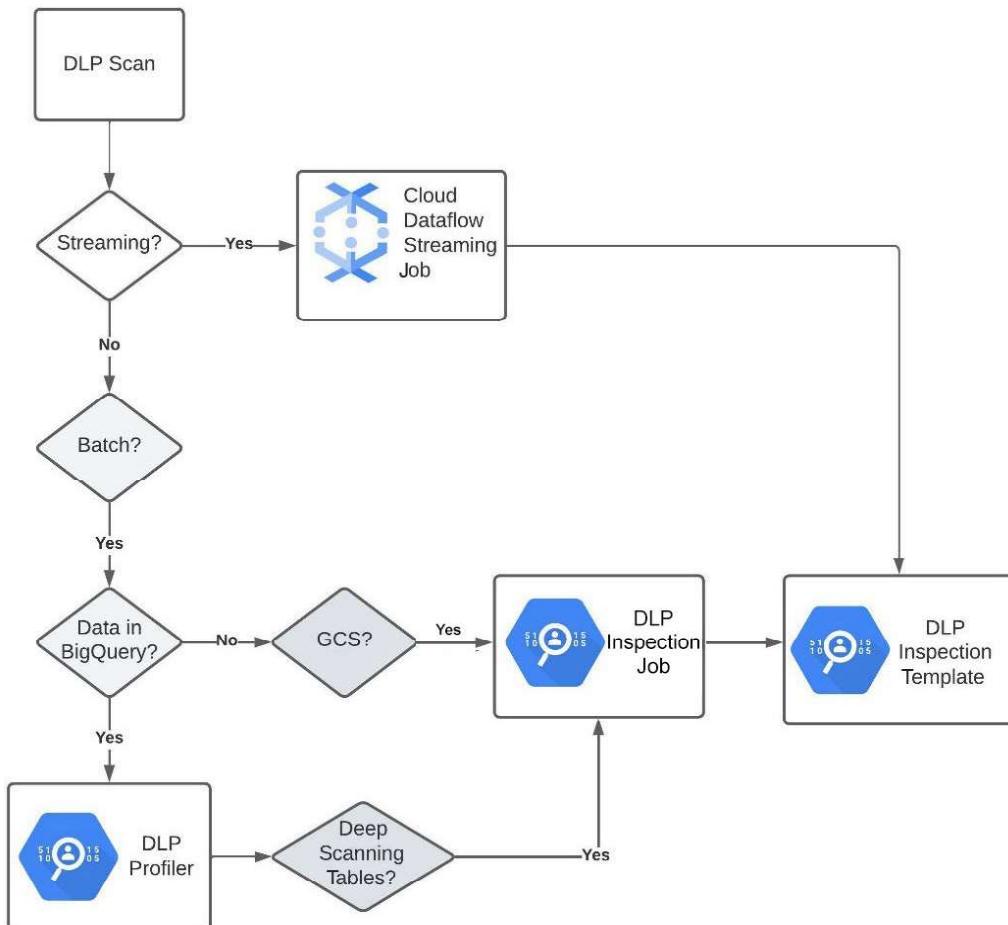


Figure 10.7 – Decision tree for inspection

Figure 10.7 shows a decision tree on when to use inspection jobs versus data profiling services.

- **Leverage built-in rules wherever possible:** Cloud DLP's built-in inspection templates should be used wherever possible. With over 150 built-in templates, Cloud DLP allows us to detect and classify data with speed and scale. For example, the US_HEALTHCARE_NPI tag can detect the 10-digit national provider number used for Medicare services. The use of built-in infoTypes saves the time and effort needed to create and maintain custom rules for standard infoTypes.
- **Trigger Cloud DLP scans based on events:** Cloud DLP's data profiling capability allows the automated scanning of all BigQuery tables and columns in your organization, folders, and projects. However, there might be a need to automate triggers that will enable the scan to run on new files that get uploaded to Cloud Storage or a new message that's added to Pub/Sub topics. To support this use case, Eventarc API or other BigQuery sink listeners can be configured to identify the change events that will eventually trigger a Cloud DLP scan.

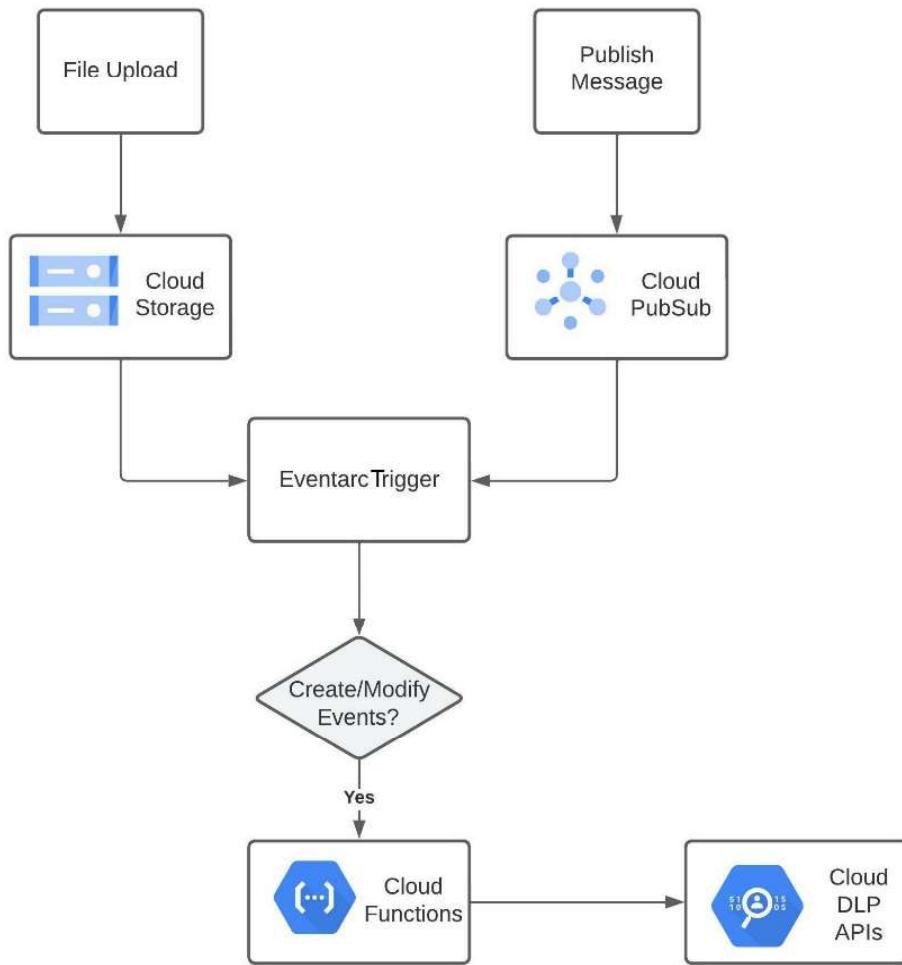


Figure 10.8 – Pattern for new updates

As shown in *Figure 10.8*, the pattern shows an example of a file upload or a Pub/Sub message triggering an Eventarc trigger. Based on the event, a cloud function will be triggered that will start the Cloud DLP scan.

- **Provide necessary permissions to the service account:** Cloud DLP creates service account agent accounts by default, which will have the necessary permissions to run Cloud DLP scan jobs. If you are scanning a highly restricted table or file, the relevant IAM roles must be added to the Cloud DLP service agent. Alternatively, if you are using a Cloud Function or Cloud Run applications to trigger Cloud DLP jobs, those service accounts must have the `dlp.jobTriggers.create` permission to successfully start the Cloud DLP scan job.
- **Publish Cloud DLP tags to Data Catalog:** Once the classification is complete, the results can be sent to Data Catalog so that the tables are tagged, only the right people get access to the data, and any sensitive data is protected. Cloud DLP allows native integration with the Data Catalog, so once we scan the BigQuery tables, it can send the tagging results into Data Catalog in the form of tag templates. The findings will be included in the summary form for the table in Data Catalog.
- **Restrict access prior to running Cloud DLP scans:** Access to any new BigQuery tables must be restricted until the Cloud DLP profiling is complete. To achieve this goal, the Cloud DLP profiler must be enabled, and only after it passes the policy checks should the data assets be moved to a state where end users can start accessing them. A quarantine state can be added to policy tags for all new tables so that users won't have access. Once Cloud DLP identifies sensitive data, a policy tag can be created in the form of a hierarchy of high, medium, and low to further restrict access depending on the content Cloud DLP identifies.

We have seen best practices for how to use Cloud DLP in various scenarios. Now let us switch gears and look at how to control data exfiltration and the best practices to prevent it.

Data exfiltration and VPC Service Controls

In the public cloud, there are several threats that organizations need to understand before deploying critical workloads. Here are a few threats that would lead to data exfiltration:

- Misconfigured IAM policies
- Malicious insiders copying data to an unauthorized destination
- Compromised code copying data to an unauthorized destination
- Access to data from unauthorized clients using a stolen credential

Here are various paths via which data can be exfiltrated in the cloud:

- Internet <-> service (stolen credentials)
 - Copy to internet
- Service <-> service (insider threat)
 - Copy from one storage service to another
- VPC <-> service (compromised VM)
 - Copy to consumer Google services
 - Copy to public GCS buckets/BigQuery dataset/GCR repo

Google Cloud offers some excellent offerings to stop the exfiltration of data as a part of its data loss prevention portfolio of products. VPC Service Controls extends a logical security perimeter around multi-tenant Google Cloud services such as **Google Cloud Storage (GCS)**, **BigQuery**, and **Google Container Registry (GCR)**. VPC Service Controls isolates Google Cloud-managed resources from the internet, unauthorized networks, and unauthorized Google Cloud resources. VPC Service Controls allows you to configure a security perimeter around the data resources of your Google-managed services (such as GCS, BigQuery, and Cloud BigTable) and control the movement of data across the perimeter boundary.

Let us look at the architecture of VPC Service Controls now.

Architecture of VPC Service Controls

Let us see a common VPC Service Controls architecture.

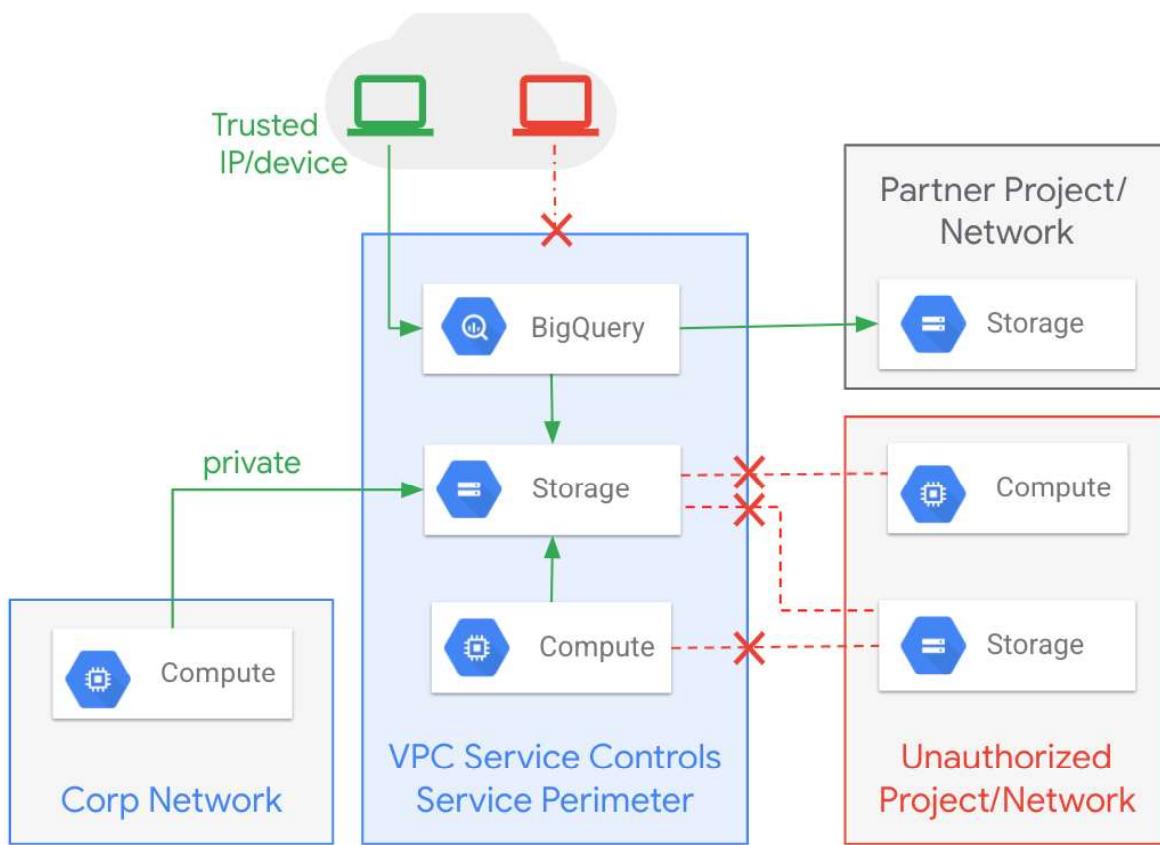


Figure 10.9 – VPC Service Controls architecture

As seen in *Figure 10.9*, the service perimeter provides a unified policy construct to isolate managed GCP resource access across the following three network interfaces:

- Internet-to-Google Cloud resource
- VPC network-to-Google Cloud resource
- Google Cloud resource-to-resource (on Google backend)

VPC Service Controls also allows you the following:

- The ability to run the configuration in dry run mode to evaluate the impact of a policy change before enforcing it.
- Fine-grained ingress and egress rules for secure data exchange between projects and other Google Cloud organizations.
- To create an access level with IP address and device attributes to enable secure access from outside the perimeter. Device attributes are typically used to establish organization trust to make sure the requests are coming from a trusted device (device attributes require BeyondCorp Enterprise).

- VPC accessible services to constrain which Google Cloud APIs (services) are accessible from a VPC network.

Now let us take a look at how VPC Service Controls can be used for API access restriction.

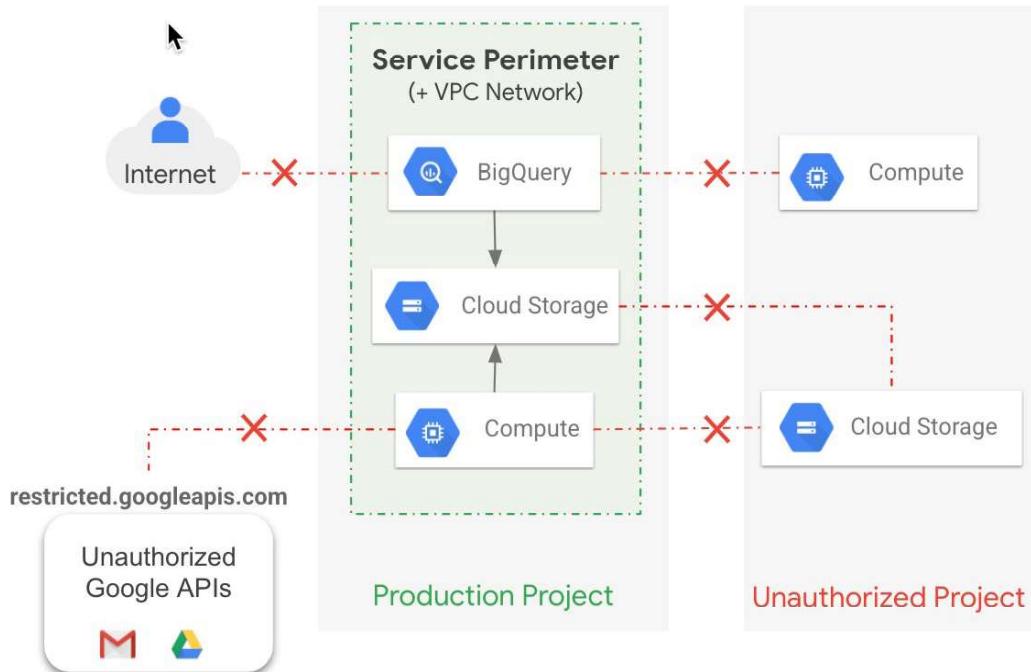


Figure 10.10 – Control APIs

As shown in *Figure 10.10*, VPC Service Controls allows you to control which APIs can be accessed from within the service perimeter:

- **Private Google Access or Private Service Connect:** Privately access Google Cloud APIs using Private Google Access.
- **restricted.googleapis.com:** Limit access to Google Cloud services that can be isolated with VPC Service Controls using the restricted **Virtual IP (VIP)** address range. To allow access to Cloud and Developer APIs protected by VPC Service Controls, use `restricted.googleapis.com`. The `restricted.googleapis.com` domain resolves to the `199.36.153.4/30` VIP range. The internet does not have access to this IP address range. The VIP `restricted.googleapis.com` denies access to Google APIs and services that are not supported by VPC Service Controls.
- **VPC accessible services:** Specify the exact list of APIs that can be accessed from a VPC.

Now we have seen how to restrict API access, let us see how we can allow access to the APIs within the perimeter.

Allowing access to protected resources within the VPC Service Controls perimeter

Use access levels to allow access to protected Google Cloud resources within service perimeters from outside a perimeter. An access level specifies a set of requirements that must be met in order for a request to be honored. Various variables, such as IP address and user identity, can be used to determine access levels.

To configure access levels, you'd use **Access Context Manager**. Let's take a look at it in more detail next.

Access Context Manager

Access Context Manager is a policy engine that allows administrators to create access control policies based on factors such as user geolocation and source IP address. These policies can be used as part of VPC Service Controls to ensure that a set of controls is applied to protect data in the project from adversaries.

Here are some rules when defining access levels:

- Only requests from outside a perimeter for resources of a protected service inside a perimeter are allowed.
- Access levels allow ingress only. Use an egress policy to allow access from a protected resource inside a perimeter to resources outside the boundary.
- The ingress and egress rules define which identities and resources are granted access to and from the perimeter. Use cases that formerly needed one or more perimeter bridges can be replaced and simplified with ingress and egress rules. A perimeter bridge is a feature of Google Cloud VPC Service Controls that allows users to securely and reliably connect two Google Cloud VPC networks. With this feature, users are able to create a secure connection between two VPC networks, allowing them to securely share resources, such as applications and databases, across the two networks. This provides an additional layer of security to protect resources from unauthorized access. Additionally, the perimeter bridge makes it easier to manage traffic between the two networks.
- Even though an access level normally accepts the external request, requests for a protected resource in a perimeter that originate from another perimeter are denied. You will need to add appropriate rules to get the desired action by applying ingress/egress rules.
- For IP-based allowlists, you can only utilize public IP address ranges in the access levels. Internal IP addresses are not allowed in these allowlists at the time of writing. A VPC network has internal IP addresses, and VPC networks must be referenced by its containing project via an entrance or egress rule, or a service boundary.

Now let us look at how to configure a VPC Service Controls perimeter.

Configuring a VPC Service Controls perimeter

To configure a VPC Service Controls perimeter, you should follow these high-level steps:

1. Set up access levels.
2. Create an access policy (aka perimeter).
3. Secure Google-managed resources with service perimeters.
4. Set up VPC-accessible services to add additional restrictions to how services can be used inside your perimeters (optional).
5. Set up private connectivity from a VPC network (optional).
6. Allow context-aware access from outside a service perimeter using ingress rules (optional).
7. Configure secure data exchange using ingress and egress rules (optional).

You also have the option to create an access level that will determine how your perimeter will be accessed. Let us look at that now.

Creating an access level

The following examples explain how to create an access level using different conditions:

1. Open the **Access Context Manager** page in the Cloud console and then open the **Access Context Manager** page.
2. If you are prompted, select your organization.
3. At the top of the **Access Context Manager** page, click **New**.
4. In the **New Access Level** pane, do the following:
 - I. In the **Access level title** box, enter a name for the access level.
 - II. Click **Add Attribute** and then select **IP Subnetworks**. The supported access level attributes are as follows:
 - i. IP subnetworks
 - ii. Regions
 - iii. Access level dependency
 - iv. Principals
 - v. Device policy
 - III. In the **IP Subnetworks** box, enter an IPv4 or IPv6 CIDR block—for example, **93.184.216.0/32**.

New Access Level

Access level title * IP_Access_Level

Access level name ? An access level name will automatically be generated based on the title.

Create conditions in Only conditions in the selected mode will be saved.

Basic mode ? Advanced mode ? Premium

Conditions

Combine conditions with
 OR AND

When condition is met, return:
 TRUE FALSE

IP Subnetworks 93.184.216.0/32

Enter one or more IPv4 or IPv6 subnetworks. Use CIDR block notation.

+ Geographic locations
+ Device policy Premium

Figure 10.11 – Creating access level

Figure 10.11 shows the New Access Level pane.

5. Click **SAVE**.
6. The access level is saved and appears in the grid on the **Access Context Manager** page.

That's it—you just created an access level that can be used in the access policy. Now let us look at how to create an access policy.

Creating a service perimeter

Access policies are always created at an organizational level. Access policies have the following components:

- **Resources you want to include in the policy:** These are folders or projects that the perimeter will protect
- **An administrator:** This is the administrative principal (a user or a service account) that either has access to view or manage the policy

- **VPC-accessible services:** This is a service that is allowed to be accessed by other VPCs
- **Access levels:** These are, as defined in the previous sections, endpoints that will have access to the perimeter
- **Ingress policy:** A rule that allows clients outside of the service perimeter to call a protected service within the perimeter
- **Egress policy:** A rule that allows services or clients within the perimeter to call a service outside of the perimeter

The screenshot shows the 'VPC Service Control Enforced Config Detail' page. At the top, there are navigation links: a back arrow, the title, and buttons for 'EDIT PERIMETER' and 'DELETE PERIMETER'. Below the title, there are several expandable sections:

- Basic details**:
 - Perimeter Title: sp_higher_trust_analytics_yzxm
 - Perimeter Name: accessPolicies/62644471578/servicePerimeters/sp_higher_trust_analytics_yzxm
 - Perimeter Type: Regular
- Projects to protect**: No projects.
- Restricted Services**: A list of Google APIs:
 - BigQuery API
 - Google Cloud Asset API
 - Google Cloud Data Catalog API
 - Google Cloud Dataflow API
 - Google Cloud Data Loss Prevention (DLP) API
 - Cloud Functions API
 - Cloud Key Management Service (KMS) API
 - Stackdriver Logging API
 - Google Cloud Pub/Sub API
 - Secret Manager API
 - Google Cloud Storage API
 - Google Compute Engine API
 - Cloud Monitoring API
 - AI Platform Notebooks API
- VPC Accessible Services**: All services allowed.
- Access Levels**: alp_higher_trust_analytic_yzxm
- Ingress policy**: No ingress policy.
- Egress policy**: (This section is partially cut off at the bottom of the screenshot.)

Figure 10.12 – Defining a service perimeter

As shown in *Figure 10.12*, the service perimeter will contain all the components necessary to create a perimeter.

Next, we will take a look at some of the best practices for VPC Service Controls.

Best practices for VPC Service Controls

Now that you understand the higher-level details of VPC Service Controls perimeters, let us go over some best practices:

- A single large perimeter is the simplest to implement and reduces the total number of moving parts requiring additional operational overhead, which helps to prevent complexity in your allowlist process.
- When data sharing is a primary use case for your organization, you can use more than one perimeter. If you produce and share lower-tier data such as de-identified patient health data, you can use a separate perimeter to facilitate sharing with outside entities.
- When possible, enable all protected services when you create a perimeter, which helps to reduce complexity and reduces potential exfiltration vectors. Make sure that there isn't a path to the private VIP from any of the VPCs in the perimeter. If you allow a network route to `private.googleapis.com`, you reduce the VPC Service Controls protection from insider data exfiltration. If you must allow access to a non-supported service, try to isolate the use of unsupported services in a separate project, or move the entire workload off the perimeter.
- An enterprise has many projects representing different workloads and applications. It is recommended that you organize these projects into folders.
- Allow ample time to gather data, conduct tests, and analyze violation logs. Make sure that stakeholders from your network, operations, and applications team are available for the task.
- Configure Private Google Access or Private Service Connect with `restricted.googleapis.com`.
- Move any applications using services not on `restricted.googleapis.com` to a separate project.
- Use dry run mode when configuring VPC Service Controls to an existing production environment. Make sure you understand the exceptions and errors produced by the dry run mode and see whether they match your design expectations. For example, certain service accounts such as an organization-level log sink will need to be allowed to collect logs. The dry run mode will catch these exceptions that you may not have considered in your design.
- Use the VPC Service Controls troubleshooting tool to analyze VPC Service Controls error codes.
- Set VPC-accessible services to restricted services to prevent access to unrestricted services.

- Ideally, do not include on-premises corp networks with user devices within a very secure service perimeter configuration.
- Document the following for every use case you use VPC Service Controls for:
 - The access pattern
 - The actors that can trigger the use case
 - Conditions that trigger the use case
 - Whether the use case is a valid access pattern and should be allowed
 - Any assumptions that pertain to the use case
 - Treat VPC Service Controls logically the same way you would a firewall; in the network, you need to “know your flows” to successfully secure the environment with the required ingress and egress flows enabled

VPC Service Controls is an important part of your cloud security strategy, so make sure you spend enough time understanding it.

Summary

We have covered a lot of things in this chapter. We went over various DLP definitions, use cases, and architecture options. We went over in detail how to use DLP for inspection and de-identification. We also saw examples of how to call DLP APIs and how to interpret the response. Last but not least, we went over data exfiltration strategies such as VPC Service Controls and their best practices. In the next chapter, we will go over the features of Secret Manager and how it can be used to store your application secrets securely.

Further reading

For more information on Cloud DLP and VPC Service Controls, refer to the following links:

- InfoType detector reference: <https://packt.link/cVC5r>
- Access level design: <https://packt.link/ZAEk4>
- Creating device-based access levels for VPC Service Controls: <https://packt.link/MpHwV>
- Ingress and egress rules for VPC Service Controls: <https://packt.link/3bBkP>

11

Secret Manager

In this chapter, we will look at Google Cloud Secret Manager. A secret is any piece of data that needs to be protected, such as passwords, encryption keys, API keys, certificate private keys, and other sensitive information. Secrets can be stored in a secure and encrypted format using Google Cloud Secret Manager. Applications that run on the cloud platform often need to be able to securely store secrets and rotate them as needed. The secrets should also provide redundancy in the event that a region goes down. Traditionally, secrets were stored in configuration files embedded within the application using some form of encryption or in clear text. However, a secure way of storing secrets is by using Secret Manager. Secret Manager is a native offering on Google Cloud to store application secrets such as database passwords or API keys. It provides a single place to store application secrets and be able to manage access and audit who did what.

In this chapter, we will cover the following topics:

- Overview of Secret Manager
- Managing secrets and versions
- Accessing a secret
- Secret rotation policy
- CMEKs for Secret Manager
- Best practices for secret management

Overview of Secret Manager

Secret Manager allows you to store and access secrets as binary blobs. It uses IAM permissions to grant access to the secrets and be able to manage them. Secret Manager is used for applications running on Google Cloud to store information such as database passwords, API keys, or certificates.

Note

Cryptographic keys should not be stored in Secret Manager. Cloud KMS is a better service since it allows you to encrypt the key material.

Before we start using Secret Manager, let us go over some core concepts. This will help you understand how Secret Manager works so you can make the right decision for your workloads.

Secret Manager concepts

Let us look at some key definitions related to Secret Manager as defined in the Google Cloud documentation:

- **Secret:** A secret is an object at the project level that stores a collection of metadata and secret versions. The metadata includes replication locations, labels, and permissions.
- **Version:** A secret version stores the actual confidential information, such as API keys, passwords, and certificates. You cannot modify a secret once you create it, but you can delete it and create a new version.
- **Rotation:** Secret Manager supports rotational schedules for secrets. This, however, doesn't create a new secret version automatically but merely sends a message to a Pub/Sub topic configured for that secret based on the rotational frequency and time.
- **Replication policy:** Secrets have global names and globally replicated metadata. The replication policy can only determine where the secret payload data is stored. Upon the creation of each secret, it will have its own replication policy. You cannot modify the location in the replication policy.
- **Resource consistency:** Adding a secret version and then immediately accessing that secret version is an example of a consistent operation in Secret Manager. Consistent operations usually synchronize in minutes but may take a few hours. You should take consistency into consideration when designing your applications.

Managing secrets and versions

Now let us look at some basic operations you can perform in Secret Manager. These operations can either be performed using `gcloud`, the Cloud console, or APIs. We have used the console to keep it simple.

Creating a secret

Follow these steps to create a secret using the Google Cloud console:

1. Go to **Console | Security | Secret Manager**.
2. Click on **Create secret**.

[←](#) Create secret

Secret details

This will create a secret with the secret value in the first version. [Learn more](#)

Name *

! Secret name is required

Secret value

Input your secret value or import it directly from a file.

Upload file [BROWSE](#)

Maximum size: 64 KiB

Secret value

Replication policy

By default, Google automatically manages where this secret is stored. If you need to manually manage this, you can customize the locations by checking the box below. All secrets are globally accessible regardless of how they are replicated and stored. The replication policy cannot be changed after a secret is created. [Learn more](#)

Manually manage locations for this secret

Encryption

This secret is encrypted with a Google-managed key by default. If you need to manage your encryption, you can use a customer-managed key instead. [Learn more](#)

Use a customer-managed encryption key (CMEK)

Figure 11.1 – Creating a new secret

As shown in *Figure 11.1*, a secret value can be copied directly into the console or provided via a file. Typically, binary secrets are provided via a file.

3. Choose the desired *rotation* period (**Set rotation period**). Note that this will only set a notification but not rotate the secret automatically.

Rotation

Setting a rotation period will send rotation notifications to Pub/Sub topics. Secret Manager will not automatically rotate the secret value. [Learn more](#)

Set rotation period

Notifications

Select Pub/Sub topic(s) that will receive event notifications whenever the secret or one of its versions is changed. These events can be user initiated changes or scheduled events. [Learn more](#)

[+ ADD TOPIC](#)

Expiration

By default, the secret never expires. To set an expiration date for this secret, select **Set expiration date** below. If you choose an expiration date, the secret will be deleted and unavailable after that time. [Learn more](#)

Set expiration date

Labels

Use labels to organize and categorize your secrets.

[+ ADD LABEL](#)

Figure 11.2 – Creating a new secret—rotation period

As shown in *Figure 11.2*, there are additional properties you can set on a secret. We recommend that you do not set **Expiration** since this can cause some unintended behavior. You can set **Labels** to organize secrets.

4. Alternatively, you can use the following `gcloud` command to create a secret:

```
gcloud secrets versions add secret-id --data-file="/file path/  
to/setcret.txt"
```

Adding a new secret version

We will now see how to add a new secret version to an existing secret:

1. Select the secret and click on **+NEW VERSION**.

Secret: "apiKey"

`projects/794301636481/secrets/apiKey`

OVERVIEW	VERSIONS	PERMISSIONS	LOGS	
Versions	+ NEW VERSION	ENABLE	DISABLE	DESTROY
	<input checked="" type="checkbox"/> Version	Status	Encryption	Created on
	<input checked="" type="checkbox"/> 1	Enabled	Google-managed	5/16/22, 5:02 PM

1 version selected

Figure 11.3 – Adding a new secret version

It should show a popup to allow you to add a secret either from a file or by copying directly into the **Secret value** field.

Add new version to "apiKey"

Input the new secret value or import it directly from a file.

Upload file

[BROWSE](#)

Maximum size: 64 KiB

Disable all past versions

[CANCEL](#) [ADD NEW VERSION](#)

Figure 11.4 – Adding a new secret version (continued)

As shown in *Figure 11.4*, you have an option to disable all past secret versions (see the **Disable all past versions** check box), so make sure to choose this option carefully as disabling a secret version currently used by your application could cause an outage.

Disabling a secret

Now let us look at how to disable a secret that you no longer want to use:

1. Click on the secret you want to disable and click **DISABLE**.

Secret: "apiKey"

projects/794301636481/secrets/apiKey

OVERVIEW		VERSIONS		PERMISSIONS		LOGS	
Versions		+ NEW VERSION		ENABLE		DISABLE	
		Version	Status	Encryption	Created on	↓	Actions
<input checked="" type="checkbox"/>	1	⚠ Disabled	Google-managed	5/16/22, 5:02 PM			⋮

1 version selected

Figure 11.5 – Disabling a secret

As seen in *Figure 11.5*, the secret will be disabled, and you will also see **Status** changed to **Disabled** on the console.

Enabling a secret

This operation will enable a secret that is in the **Disabled** state:

1. Click on the secret you want to enable and click **ENABLE**.

Secret: "apiKey"

projects/794301636481/secrets/apiKey

OVERVIEW		VERSIONS		PERMISSIONS		LOGS	
Versions		+ NEW VERSION		ENABLE		DISABLE	
		Version	Status	Encryption	Created on	↓	Actions
<input checked="" type="checkbox"/>	1	✓ Enabled	Google-managed	5/16/22, 5:02 PM			⋮

1 version selected

Figure 11.6 – Enabling a secret

As seen in *Figure 11.6*, the secret will be enabled, and you will also see **Status** changed to **Enabled** on the console.

Accessing a secret

Accessing a secret version returns the secret contents as well as additional metadata about the secret version. When you access a secret version, you specify its `version-id`. You can also access the latest version of a secret by specifying `latest` as the version.

Accessing a secret version requires the `secretmanager.secretAccessor` IAM role. Typically, this role is granted to the service account used by your application.

The following is a `gcloud` command that can be used to access a particular version of the secret. However, a common method is to use application libraries for access, as you will see later in the section:

```
gcloud secrets versions access version-id --secret="secret-id"
```

You can also use the `latest` keyword to get the current version, but this is *not* a recommended best practice.

Accessing a binary secret version

You can access binary secrets directly but note that Cloud SDK formats the output as UTF-8, which can corrupt binary secrets. To get the raw bytes, Cloud SDK prints the response as base64-encoded, which you can then decode:

```
gcloud secrets versions access version-id --secret="secret-id"  
--format='get(payload.data)' | tr '_-' '/+' | base64 -d
```

Accessing secrets from your application

You can access secrets from Google Cloud services such as Compute Engine, Application Engine, GKE, and Cloud Run.

Secret Manager has a REST and gRPC API for using and managing secrets. However, a convenient way to get access to secrets is by using client libraries. You can use these client libraries bundled with your application to access secrets. You do not need to embed service account keys within your application to access Secret Manager APIs. All you need is the right IAM role for the service account used by the Google Cloud service of your choice, such as GCE, GAE, or Cloud Functions:

- In Cloud Build, you can access secrets using environment variables in the build step of the build YAML config file by referring a field called `availableSecrets` to a specific secret version.
- For Cloud Run, you can use Secret Manager in two ways:
 - Mount the secret as a volume, which will make the secret available to the container as files. Reading the volume will read the secret from Secret Manager. You can use the `latest` keyword to fetch the secret using this method.

- Access the secret using a traditional environment variable. Environment variables are resolved by Cloud Run at the start of the container. When you use this method, do not use the `latest` keyword, use a particular version.
- For GKE, you can use Workload Identity in conjunction with Secret Manager client libraries to access secrets. You can also use the Secret Store CSI driver in conjunction with the Google Secret Manager provider for this driver. This allows you to access secrets as files mounted in Kubernetes pods. The provider is an open source project supported by Google that uses the workload identity of the pod that a secret is mounted onto when authenticating to the Google Secret Manager API. For this to work, the workload identity of the pod must be configured and appropriate IAM bindings must be applied. This plugin is built to ensure compatibility between Secret Manager and Kubernetes workloads that need to load secrets from the filesystem. The plugin also enables syncing secrets to Kubernetes-native secrets for consumption as environment variables. When evaluating this plugin, consider the following threats:
 - When a secret is accessible on the filesystem, application vulnerabilities such as **directory traversal** (`https://packt.link/wYWmj`) attacks can become higher severity as the attacker may gain the ability to read the secret material.
 - When a secret is consumed through environment variables, misconfigurations such as enabling a debug endpoint or including dependencies that log process environment details may leak secrets.
 - When syncing secret material to another data store (such as Kubernetes Secrets), consider whether the access controls on that data store are sufficiently narrow in scope.
- For multi-cloud application deployments, consider using workload identity federation on other platforms, such as Amazon Web Services or Microsoft Azure, that use current identity protocols to authenticate to Google Cloud APIs. Typically, you would want to use the native offering of the cloud providers for the secrets where your application is deployed.

Let us look at how you can achieve multi-region redundancy of secrets by using a secret rotation policy. This strategy supports the application redundancy you would need for high availability and failover.

Secret replication policy

Secrets are a global resource entity; however, secret payloads (the underlying secret material) are stored locally within a region. Some regulated customers such as financial and healthcare institutions may have strict regionalization requirements, while other customers may want to store the secret near the data. A replication policy allows control over where secret payloads are stored.

There are two replication policy types: automatic and user-managed.

Automatic

With the automatic policy type, the replication of the secret is managed by Google. This policy provides the highest level of availability:

- When a secret has an automatic replication policy, its payload data is copied as many times as needed. This is the easiest way to set things up, and most users should choose it. This is the policy that is used by default when a secret is created using the Google Cloud CLI or the web UI.
- A secret that is automatically replicated is stored in a single place for billing purposes.
- For a resource location organization policy evaluation, you can only make a secret with an automatic replication policy if you can make resources global.

User-managed (user-selected)

With the user-managed replication policy type, the secret is replicated to only locations selected by users:

Note

The secret version addition will be impacted if one of the regions you selected is unavailable.

- The payload data of a secret with a user-managed replication policy is duplicated to a user-configured list of locations. The secret can be duplicated to as many supported locations as you choose. If there are regulations on where the secret payload data can be stored, this could be useful.
- Each location in the user-managed replication policy is treated as a separate location for billing reasons.
- A secret with a user-managed replication policy can only be created if the organization policy allows the creation of resources in all selected locations.
- As a rule, choose between two and five locations for replication. Having a large number of replication locations results in longer replication delays as well as increased latency in accessing the new secret payload.

Let us look at some of the encryption options, namely Google default encryption and customer-managed encryption, for Secret Manager.

CMEKs for Secret Manager

By default, secrets are encrypted with Google default encryption. However, some highly regulated customers require control of keys, so Secret Manager supports **customer-managed encryption keys (CMEKs)** (within Cloud KMS) for encrypting:

Note

However, if you disable or permanently destroy the CMEK, the secret encrypted with that key cannot be decrypted.

- Secret payloads are encrypted by Google-managed keys before being written to persistent storage with no additional configuration required.
- Secret Manager encrypts data with a unique **data encryption key (DEK)** before writing it to persistent storage in a specific location. The Secret Manager service owns a replica-specific key called a **key encryption key (KEK)**, which is used to encrypt the DEK. This is commonly referred to as envelope encryption.
- The CMEK is a symmetric key that you control within Cloud KMS when using CMEKs with Secret Manager. The CMEK must be stored in the same GCP region as the secret. This makes sure that the Secret Manager has access to the key in the event of a zone or region failure.
- Customers who desire complete control over encryption keys can use the **Cloud External Key Manager (EKM)** key in the CMEK policy for encryption and decryption. This completely removes Google's management of the encryption key for your secrets as the encryption key is controlled outside of Google Cloud.
- CMEKs are also supported when using user-managed replication. When you create a CMEK in Cloud KMS, make sure the location of the key matches the location of the secret if you are using a user-managed replication policy.

So far, we have seen how to set up secrets in Google Cloud. Let us now go over some of the best practices for managing secrets in production and development environments.

Best practices for secret management

Here are some general best practices that Google recommends when it comes to managing secrets:

- As with pretty much all services in Google Cloud, access to the Secret Manager API is protected by IAM. Follow the principle of least privilege when granting permissions to secrets to your applications.

- Divide applications and environments (staging/production) into independent projects. This can assist in segregating environments with IAM binding at the project level and guarantee that quotas are implemented independently.
- If necessary, establish a custom role or choose an existing role with the bare minimum of access. Think about who manages the secret (creates the secret, disables/enables it, or creates a new version) and who uses it, such as developers. You should have a separation of duties between these two roles, especially in production.
- Use secret-level IAM bindings or IAM conditions to limit access to the necessary subset of secrets when numerous services' secrets are in a single project.
- Use `gcloud auth application-default login` when developing locally. This produces a file with credentials that client libraries will immediately recognize.

Let us look at some of the best practices that you should consider while specifically developing your application.

Best practices for development

Here are some best practices for working with Secret Manager during the development of an application that will ultimately be deployed in Google Cloud:

- To use secrets within your development IDEs such as Visual Studio Code or IntelliJ, use the Secret Manager integration plugin. Please see the Google Cloud documentation on how to enable Secret Manager for your choice of IDE (<https://packt.link/XRAGF>).
- For deployments in a single project, consider creating secrets in the project where the workload will be deployed. This makes the management of secrets much easier.
- Do not use slashes, /, when naming a secret as this violates Google Cloud's resource naming convention.
- Secret data stored in secret versions must fit within the quotas and limits specified in the Secret Manager specification. Apart from this, Secret Manager has no restrictions on the format of the secret material.
- Secrets such as API keys, SSH keys, and cryptographic keys are generated in external systems with formatting requirements.
- Use randomly generated, large, complicated values and cryptographically strong random number generators for secrets such as passwords when the format and length are not constrained.

Best practices for deployment

Here are some best practices for Secrets Manager while deploying applications on Google Cloud:

- Use a separate project for secrets that are used by workloads in many projects, separate from the workload projects.
- Instead of granting Cloud IAM roles at the organization, folder, or project level, grant them directly on secrets. This will prevent over-grants and misuse.
- Managing secret-level IAM role grants can be easier using an **infrastructure-as-code (IaC)** approach such as Terraform. You can typically embed this in the given application CI/CD pipeline.
- The expiration date feature shouldn't be used for production workloads. This feature is best used to clean up transitory situations automatically. Consider using time-based IAM conditional access to secrets instead of secrets that expire.
- On projects that contain secrets or enclosing folders or organizations, avoid utilizing basic project-wide roles such as owner, editor, and viewer.
- Create distinct key administration and access roles according to the following guidelines:
 - Most applications require secret versions for specific secrets. For those specific secrets, grant these accounts `roles/secretmanager.secretAccessor`.
 - Grant `roles/secretmanager.secretVersionAdder` to processes that add new secret versions, for example, a CI/CD pipeline or administrator that owns the secret.
 - Grant `roles/secretmanager.secretVersionManager` to processes that handle timed disabling and destruction of secret versions.
- Rotate your secrets periodically to minimize the impact of leaked secrets.
- Exercise the rotation flow on a regular basis to lessen the chances of an outage.
- Unless your workload has specified location requirements, use the automatic replication policy. Most workloads' availability and performance requirements are met by the automated policy.
- To obtain and analyze `AccessSecretVersion` requests, enable data access logs at the folder or organization level.
- Automatically rotate secrets and have emergency rotation processes ready in case of a compromise.
- Instead of utilizing the most recent alias, refer to secrets by their version number. Create a hidden version of your program that is read at launch. Use your existing release methods to deploy updates to version numbers. Although using the most recent alias can be easy, if the latest version of the secret has an issue, your workload may be unable to use the secret version. The configuration can be evaluated and rolled back using your regular release processes if you pin it to a version number.

- Before deleting or destroying secret versions, disable them. This helps prevent outages by putting the secret in a reversible state that looks like destruct. That is, before permanently deleting data, you can disable it and wait a week to ensure there are no lingering dependencies.
- Secret Manager should not be used to store Google Cloud service account keys.
- In addition to IAM controls, you can set up a VPC Service Controls perimeter for your organization to limit access to the Secret Manager API from outside of Google Cloud using network-based controls such as IP address range and service accounts.
- The organization policy `PolicyMemberDomains` can be used to limit the identities that can be added to IAM policies for secrets.
- Estimate your peak secret usage (taking into account a *thundering herd* of queries owing to concurrent application deployments or service autoscaling) and make sure your project has enough quota to accommodate it. Request an increase in quota if more is required.
- Using Cloud Asset Inventory, keep track of secrets across your business to do the following:
 - Assist in the discovery of secrets throughout your company.
 - Identify non-conformance with organizational standards such as rotation, encryption setup, and location.

Having discussed the best practices for production, we now turn our attention to troubleshooting and auditing Secret Manager operations.

Secret Manager logs

You can consult Cloud Logging to audit or troubleshoot the operation of Secret Manager.

Let us assume `project_id = acme-project-id`, `folder_id = acme-folder`, `billing_account_id = 123456`, and `organization_id = 987654321`. The logs of interest would have the names listed as follows:

```
projects/acme-project-id/logs/cloudaudit.googleapis.com%2Factivity
projects/acme-project-id/logs/cloudaudit.googleapis.com%2Fdata_access
projects/acme-project-id/logs /cloudaudit.googleapis.com%2Fsystem_event
projects/acme-project-id/logs/cloudaudit.googleapis.com%2Fpolicy

folders/acme-folder/logs/cloudaudit.googleapis.com%2Factivity
folders/acme-folder/logs/cloudaudit.googleapis.com%2Fdata_access
folders/acme-folder/logs/cloudaudit.googleapis.com%2Fsystem_event
folders/acme-folder/logs/cloudaudit.googleapis.com%2Fpolicy

billingAccounts/987654321/logs/cloudaudit.googleapis.com%2Factivity
```

```
billingAccounts/987654321/logs/cloudaudit.googleapis.com%2Fdata_
access
billingAccounts/987654321/logs/cloudaudit.googleapis.com%2Fsystem_
event
billingAccounts/987654321/logs/cloudaudit.googleapis.com%2Fpolicy

organizations/987654321/logs/cloudaudit.googleapis.com%2Factivity
organizations/987654321/logs/cloudaudit.googleapis.com%2Fdata_
access
organizations/987654321/logs/cloudaudit.googleapis.com%2Fsystem_
event
organizations/987654321/logs/cloudaudit.googleapis.com%2Fpolicy
```

This concludes the logging part of the Secret Manager application.

Summary

In this chapter, we reviewed Secret Manager, its operations, and some critical aspects of setting secrets for your requirements. We also discussed best practices for the development and deployment of secrets. It is important to follow these best practices so that your application is designed optimally and to reduce the risk of outages. As a member of the application development team or information security team, you should be able to design your applications to leverage Secret Manager.

In the next chapter, we will look at the logging features of Google Cloud. Cloud Logging is one of the critical aspects of cloud deployment. We will also go over the best practices of Cloud Logging.

Further reading

- Secret Manager IAM roles: <https://packt.link/wgzCw>
- Secret Manager rotation: <https://packt.link/slcoa>
- Creating and managing expiring secrets: <https://packt.link/js5eW>
- Analyzing secrets with Cloud Asset Inventory: <https://packt.link/v7w8z>
- Secret Manager data integrity assurance: <https://packt.link/8AmOG>

12

Cloud Logging

Logging provides visibility into your environment and aids in troubleshooting and incident response. In this chapter, we will discuss what Cloud Logging is, how it works, the different types of logs and their applications, and how to collect, store, analyze, and export logs. We will also look at who log producers and consumers are and how to export logs to a centralized logging solution or a **Security Information and Event Management (SIEM)** system running either on-premises or in the cloud. Finally, we will discuss how to securely store and keep logs to meet regulatory requirements.

In this chapter, we will cover the following topics:

- Overview of Google Cloud Logging
- Understanding log categories
- Log management
- Logging and auditing best practices

Introduction to Google Cloud logging

Cloud Logging is a managed service on Google Cloud. It gives you the ability to collect, store, search, analyze, monitor, and set alerts based on logs that are collected. You can use Cloud Logging to collect log data from over 150 applications, Google Cloud components, third-party cloud providers, and any combination of cloud platforms (hybrid environments). Cloud Logging, formerly known as Stackdriver, is part of Google Cloud's operations suite. It includes a console interface called the Logs Explorer, query logs, and an API to manage logs programmatically.

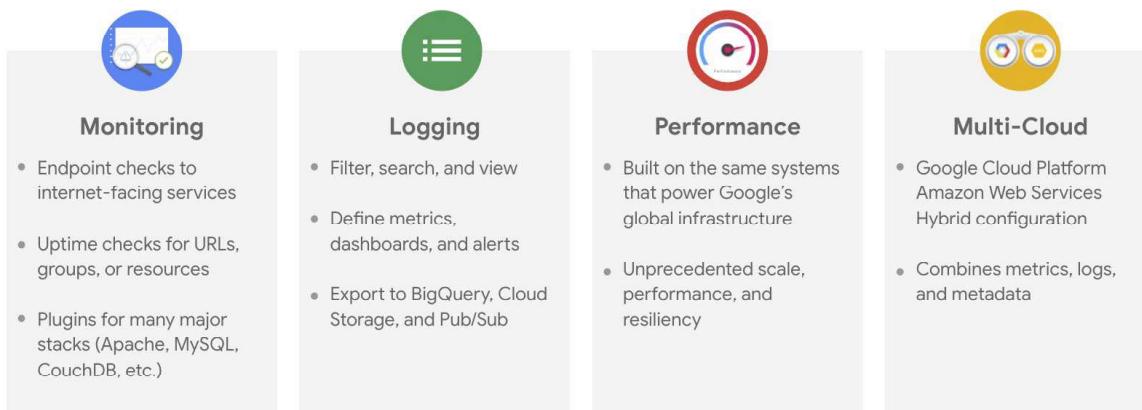


Figure 12.1 – Google Cloud's operations suite

Cloud Logging is one of the four components of Google Cloud's operations suite, as shown in *Figure 12.1*. Let's explore each of these components in more depth. Cloud Logging captures all logs and provides searching and filtering capabilities. Cloud Logging derives metrics from logs, creates dashboards, autoscales instances, and provides the ability to export logs for retention and analytics. We will cover these topics in more depth later in this chapter.

Application Performance Management (APM) combines the monitoring and troubleshooting capabilities of Cloud Logging and Cloud Monitoring with Cloud Trace, Cloud Debugger, and Cloud Profiler, to help you reduce latency and cost so you can run more efficient applications. This is out of the scope of the Google Cloud Professional Security Engineer exam.

Additionally, the Google Cloud operations suite offers multi-cloud capabilities in its support of both Google Cloud and Amazon Web Services and can also monitor on-premises resources and integrate with partner solutions.

By default, the operations suite monitors all system metrics, including Google **Compute Engine** (GCE) and **Google Kubernetes Engine** (GKE). The operations suite also monitors applications and managed services, such as Datastore and BigQuery. You can utilize the operations suite's Monitoring API to interact with users and collect business metrics. You can also deploy a monitoring agent on compute instances to collect lower-level and third-party tool metrics, such as disk latency/IO, for supported services. *Figure 12.2* illustrates the capabilities of the operations suite and how cloud resources generate signals and telemetry that are then consumed by different operations suite functions.

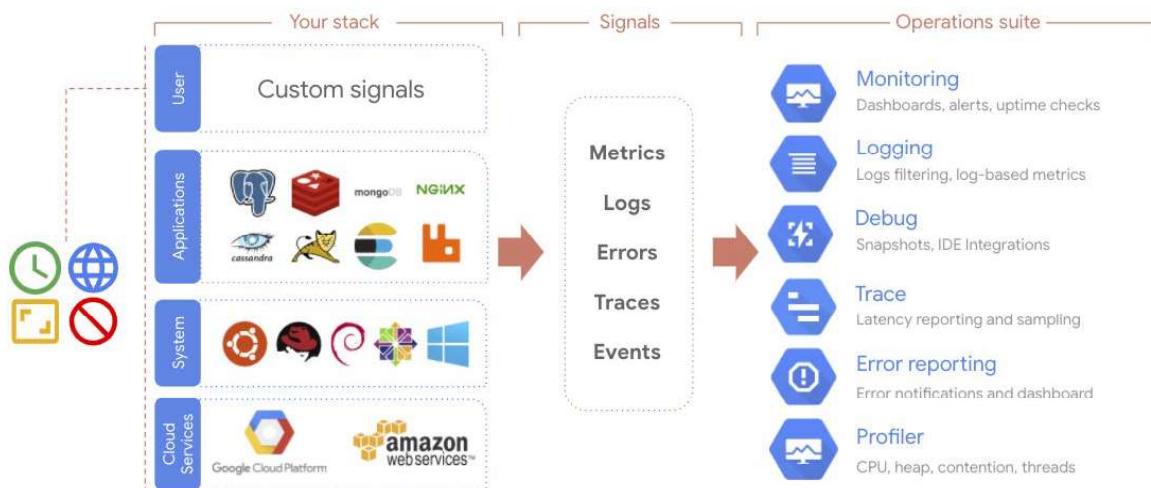


Figure 12.2 – Operations suite layers

We will now go over some fundamental Cloud Logging principles, as well as a few significant terms. It is important to understand them as they will be used throughout the chapter and assessed in the exam:

- **Log entries:** A log entry is a record of the state of a computer system or a description of specific events or transactions. Each log entry includes different types of data based on the service generating the log. Most log entries contain these attributes:
 - A timestamp to indicate when the event took place or when it was received by Cloud Logging.
 - The resource that produced the log entry.
 - A payload, also known as a message, is either provided as unstructured textual data or as structured textual data in JSON format.
 - The name of the log to which it belongs. Some logs also have metadata associated with them, such as the severity of each associated log entry.

- **Logs:** In Google Cloud resources, such as a Google Cloud project, a log is a named collection of log entries. There can be no logs unless there are log entries. A log's name is identified by the full path of the resource to which the log entries belong, followed by a simple log ID, such as `syslog`, or a structured ID that includes the log's writer, such as `compute.googleapis.com/activity`.
- **Retention period:** Log entries are stored in Cloud Logging for a specific period of time known as the retention period. Each log type has a different default retention period. You have the ability to create custom retention periods to store logs for longer periods than the defaults that are available in Cloud Logging.
- **Queries and filters:** Queries are written in the Logging query language. Queries return the matching log entries. The Logs Explorer and the Logging API both employ queries to select and display log entries.
Filters are queries in the Logging query language that are used by log sinks to direct logs that match a given expression to a storage destination. Filters are also used to redirect matching logs to Cloud Monitoring.
- **Log Router:** In order to reach the Cloud Logging API, all collected logs must first go through the Log Router. For each log entry, the Log Router checks against pre-existing rules to identify which log entries should be ingested (stored) and routed to a destination, as well as which log entries should be excluded (discarded). We will cover the Log Router in more detail later in this chapter.
- **Sinks:** A sink consists of a filter and a storage destination. The filter logs entries to send to the sink.
- **Log-based metrics:** You can define metrics for the logs and then create alerts and triggers based on those metrics.
- **Access control:** You can control who has access to the logs using Cloud IAM permissions.

Each of the log types is categorized based on the log type, such as security, user, or platform. We will look at all these categories in more detail in the next section.

Log categories

In this section, we broadly divide logs into three categories. We will look at each category from the perspective of security.



Figure 12.3 – Log categories

Figure 12.3 illustrates the different types of log categories: security logs, user logs, and platform logs. Security logs consist of admin activity logs, data access logs, system event logs, and transparency logs. User logs are generated by user software, services, or applications and are written to Cloud Logging using a logging agent, the Cloud Logging API, or the Cloud Logging client libraries. Google Cloud Platform logs are service-specific logs that can help you better understand the Google Cloud services you're using. VPC flow logs, firewall logs, and other API logs are examples of platform logs.

Our focus will be only on logs that are useful from a security perspective. That doesn't necessarily mean that we will only look at the security logs category. We will also look at user logs and platform logs, specifically the subordinate log types within these categories that pertain to security.

Security logs

Cloud Logging offers two forms of security logs: **Cloud Audit Logs** and **Access Transparency Logs**. Cloud Audit Logs generates the following types of audit logs for each cloud project, folder, and organization:

- **Admin Activity audit logs** include information on API calls and actions that potentially alter cloud resource settings or metadata. Activities such as creating a virtual machine or altering Cloud IAM permissions are examples of such changes that are captured by Admin Activity audit logs. You can't configure, exclude, or disable these types of logs. They're always recorded for security reasons. Admin Activity audit logs are created even if the Cloud Logging API is disabled. The retention period for administrator activity records is 400 days.
- **Data Access audit logs** show who accessed which data, the access time, and what actions they performed. Data Access audit logs are retained for 30 days. Not all Google Cloud services are supported. With the exception of BigQuery, Data Access audit logs are disabled by default.
- **System Event audit logs:** Events that alter resource configurations are recorded in System Event audit logs. These logs are generated automatically by Google's systems, not by users. System Event audit logs cannot be configured, excluded, or disabled; they are always written.
- **Policy Denied audit logs:** Policy Denied audit logs are created when a Google Cloud service rejects a user or service account access for a security policy violation. VPC Service Controls determines security policies and sends Cloud Logging Policy Denied audit logs. Your Google Cloud project gets charged for Policy Denied audit log storage. You can't disable Policy Denied audit logs, but you can exclude them from Cloud Logging.

The next type of security log is **Access Transparency**. Access Transparency provides logs based on the actions taken by Google staff when accessing your Google Cloud content. Access Transparency logs can help you track the compliance posture of your organization with regard to your legal and regulatory requirements. We covered Access Transparency in depth in *Chapter 3, Trust and Compliance*.

User logs

As mentioned earlier, user logs, sometimes referred to as user-written logs, are written to Cloud Logging using an agent, the Cloud Logging API, or client libraries. Let's take a look at the user-agent-based log type:

- **Cloud Logging agent:** An agent can be installed to stream logs from instances to Cloud Logging. The Cloud Logging agent is based on Fluentd and is preconfigured by default to send syslogd logs and common third-party tools' logs (for example, MySQL, PostgreSQL, MongoDB, Cassandra, and Apache). It is possible to customize the pre-configured Fluentd agent configurations as required.

Platform logs

Next, we will look at examples of platform logs:

- **VPC flow logs:** VPC flow logs contain granular VM flow-level network telemetry from Google Cloud. VPC flow logs provide insights into VPC traffic and can help with performance, traffic planning, a better understanding of networking costs, and forensics in the case of security-related incidents. VPC flow logs are aggregated collections of packets for a given interval and for a given connection. These logs are by default retained for 30 days and they do not lead to any performance degradation. The VPC flow log format includes a 5-tuple—for the source IP address, destination IP address, source port number, destination port number, and protocol (TCP or UDP)—and the timestamp. The metrics include packets, bytes (throughput), RTT for TCP flows, VPC annotations (region, zone, and VM name), and geographic annotations (country, region, and city). Logs are collected on each VM connection and aggregated at 5-second intervals. Google samples about one out of every 10 packets.
- **Firewall logs:** You can enable firewall rule logging individually for each firewall rule. Firewall rule logging only records TCP and UDP connections. Log entries are written from the perspective of VM instances.
- **NAT logs:** You can enable NAT logging on a NAT gateway instance. You can choose to generate logs for the creation of NAT connections and when packets are dropped because no port was available for NAT. Cloud NAT logs both TCP and UDP, including dropped egress packets. Dropped ingress packets (for example, if an inbound reply to an egress request was dropped) are not logged.

We have covered the different types of logs and their relevant examples. We will now consider the retention of logs. Defining a retention period gives you the ability to change the log period to match your compliance or regulatory requirement. Let's take a look at some of the key log types and their associated log retention periods in the following section.

Log retention

Cloud Logging retention is critical for many organizations due to compliance requirements. The default retention period is often not sufficient; therefore, many customers configure custom retention periods for their logs. Longer retention periods come at a higher cost.

	Retention	Cost
Admin Activity audit logs	13 months (400 days)	Free
System Event audit logs	13 months (400 days)	Free
Access Transparency logs	13 months (400 days)	Free
Data Access audit logs	30 days	\$0.50 per GB over 50GB/Month
All other logs	30 days	\$0.50 per GB over 50GB/Month
Workspace admin console logs	6 months	Free
Log sink	Indefinite	Based on service option (GCS, BQ, Pub/Sub)

Figure 12.4 – Cloud Logging default retention periods and costs

Figure 12.4 shows the default retention periods and the associated costs for the different types of logs that we have covered. The costs are correct at the time of writing and may change.

Next, we'll explore the topic of log management. We will learn about log producers and consumers and understand how a log router works. We will also discuss exporting logs and creating log sinks in detail.

Log management

A key component of logging is the ability to export logs so they can be consumed by a service running on the cloud or by a third-party solution. You will also learn how to aggregate logs, build log pipelines, perform log analysis, and deliver logs to SIEM solutions.

Log producers

All resources that generate logs, such as products, services, and applications, whether they are in the cloud or a hybrid environment, are classified as log producers. These could be Google Cloud services or services running in your on-premises data center or a third-party cloud service provider.

Log consumers

Log consumers are services such as centralized logging solutions and SIEM solutions. Any service that has the ability to consume logs and provide functions such as alerting, troubleshooting applications, and business intelligence is considered a log consumer. Cloud Logging on Google Cloud is an example of such a service.

Log Router

The term Log Router has been used a few times in the chapter; we will now look at the workings of the Log Router in more detail to understand what it is and how it works. A Log Router ensures that logs are collected from producers and delivered to the consumer.

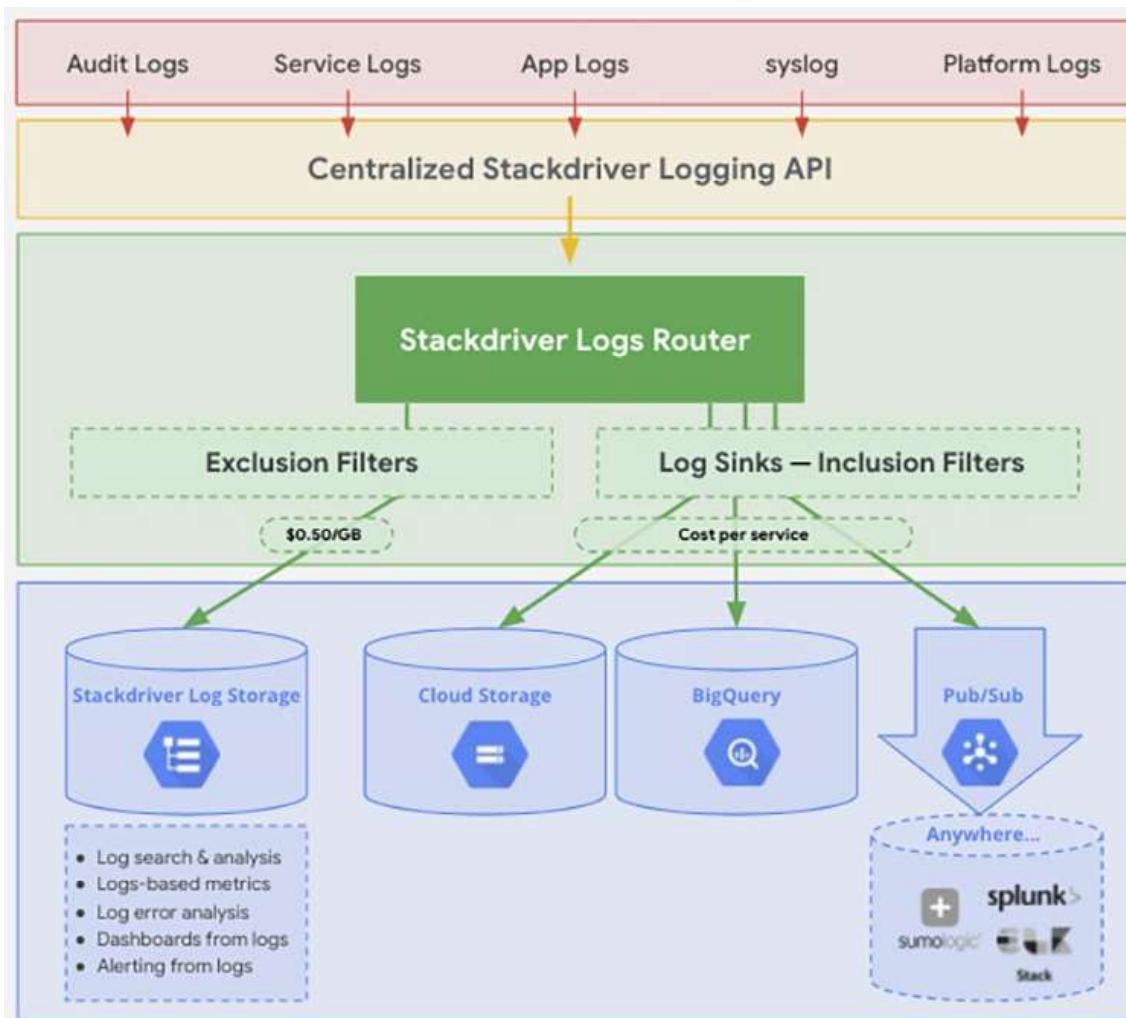


Figure 12.5 – Log Router

Figure 12.5 shows how a Log Router works. Logs are transmitted to the Cloud Logging API, where they are processed by the Log Router. These logs include audit logs, platform logs, and user-written logs. The Log Router compares each log entry to the rules that are already in place to determine which log entries should be ingested and stored, which log entries should be included in exports, and which log entries should be discarded.

To export logs, create a filter that selects the relevant log entries to export, and then select a destination. You can choose from these options:

- JSON files can be stored in Google Cloud Storage buckets.
- Tables can be created in BigQuery datasets.
- JSON messages can also be delivered via Pub/Sub topics. This option also supports third-party integrations, such as Splunk.

Alternatively, you can create another Google Cloud project and have the log entries stored in Cloud Logging log buckets there.

Sinks are used to store the filter and the final destination. Google Cloud projects, organizations, files, and billing accounts can all be used as sinks. We covered sinks earlier in this chapter. We will look at them in more detail when we look at log exports in the following section.

Log sinks and exports

Before we look at exporting logs, it's important to understand log sinks and inclusion/exclusion filters. Sinks tell Cloud Logging where to send logs. With sinks, you can send some or all of your logs to destinations that can handle them. Sinks are part of a specific Google Cloud resource, such as a project, billing account, folder, or organization. When a resource gets a log entry, the Log Router sends the log entry to the appropriate sinks and, if enabled, the Log Router forwards it to any ancestral sinks that belong to the resource hierarchy. The log entry is ultimately sent to the destination associated with the sink.

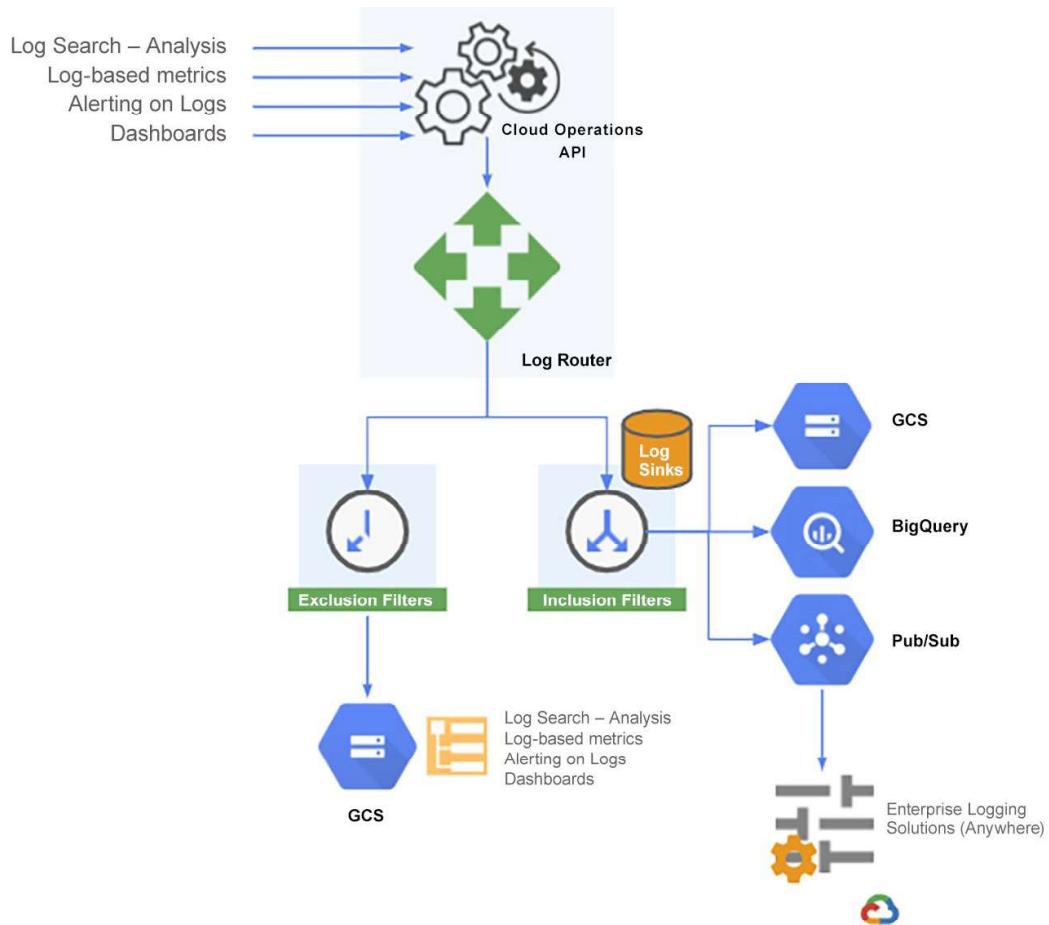


Figure 12.6 – Log exports and sinks

Figure 12.6 shows the relationships between sinks, filters, and exports. By defining one or more sinks that have a log query and an export destination, you can easily export logs from a database. New log entries are compared to each sink as they are received by Cloud Logging. When an entry matches a sink's filter, the entry is written to the sink's destination.

If no filters are attached to a sink, all logs will match and be forwarded to their final destination. The inclusion filter can be used to choose specific logs for the sink. Alternatively, you can use exclusion filters to keep unwanted logs out of the sink altogether. Filters are defined using the Logging query language.

Log buckets are containers used by Cloud Logging to store and organize log data in your Google Cloud projects, billing accounts, folders, and groups. Using Cloud Logging, you can study your logs in real time by having them indexed, optimized, and provided to you. Cloud Logging buckets are distinct from Cloud Storage buckets, despite the similarity of their names. Log buckets should be considered a local resource. You can pinpoint the location of the servers and storage facilities that handle your logs.

Next, we will look at the steps for how to create a log export.

Creating a log export

In this section, you will learn how to create a log export sink. You will build a centralized logging solution by creating log buckets that are bound to specified regions with retention policies.

You can create 100 log buckets per project. Follow these steps to create user-defined log buckets:

1. From the **Logging** menu, select **Logs Storage**.
2. Click **Create Logs Bucket**.
3. Fill out **Name** and **Description** for your bucket.
4. (*Optional*) To set the region in which you want your logs to be stored, click the **Select log bucket region** drop-down menu and select a region. If you don't select a region, the global region is used, which means that the logs could be physically located in any of the regions. In our demonstration, in *Figure 12.6*, we have used a region in Australia.

Note

You can't change the region of your log bucket once you've created it. You can, however, create an additional log bucket in a different region, reroute the associated sinks to the new log bucket, and then delete the log bucket.

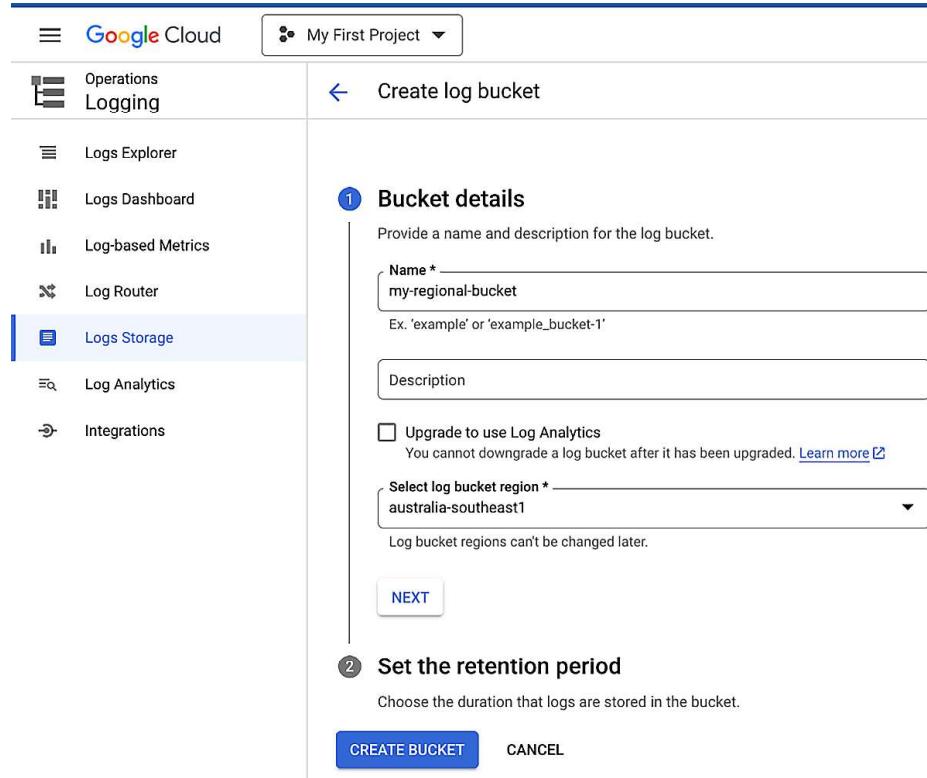


Figure 12.7 – Create log bucket

Figure 12.7 shows how to configure a log bucket. Once you have specified the name of the bucket as per the naming convention, you can then either set the bucket region to global or select the region of your choice.

5. (*Optional*) To set a custom retention period for the logs in the bucket, click **NEXT**.
6. In the **Retention period** field, enter the number of days that you want Cloud Logging to retain your logs. Google Cloud lets you define that period in seconds, days, and years, with a maximum period of retention of 100 years. If you don't customize the retention period, the default is 30 days. In *Figure 12.8*, we have kept the retention period set to its default value. You can also update your bucket to apply custom retention after you create it. Keep in mind that longer retention periods will result in higher charges.

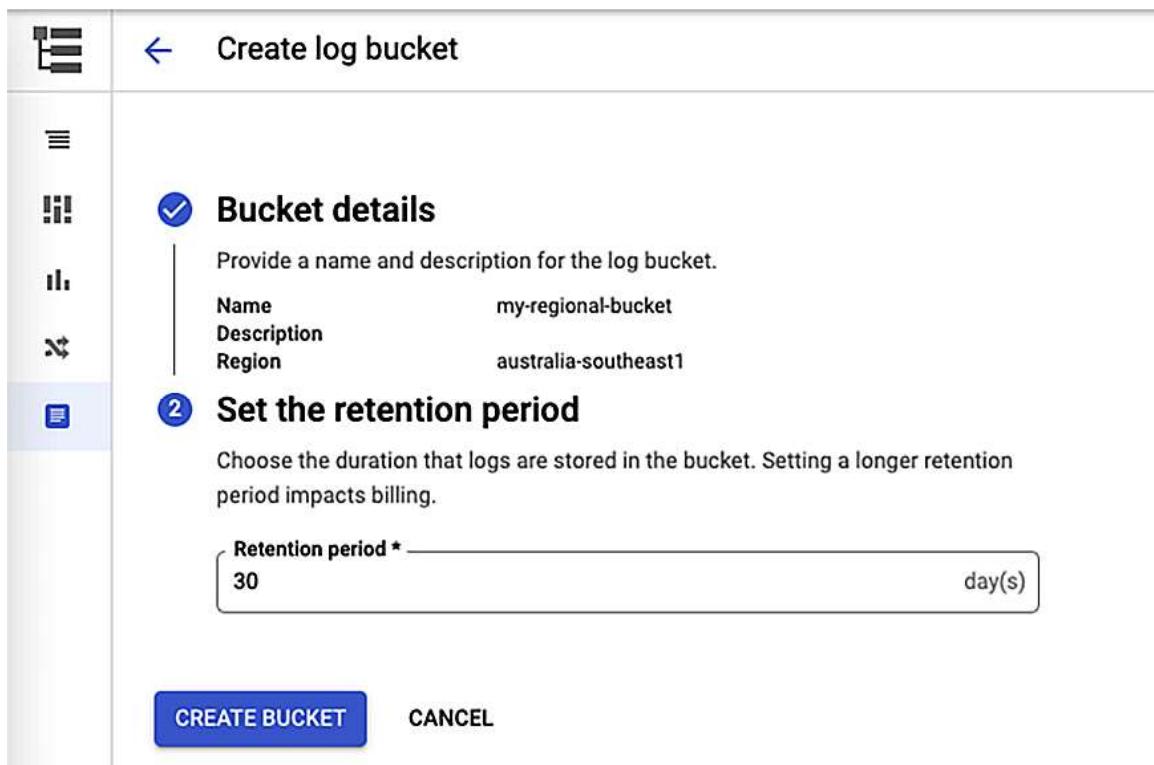


Figure 12.8 – Set the retention period for the log bucket

7. Click **CREATE BUCKET**. Your new bucket will appear in the **Log bucket** list.

You can also lock the bucket. When you stop updates to a bucket, you also stop updates to the bucket's retention policy. After you lock a retention policy, you can't delete a bucket until every log in it has been kept for as long as the policy dictates.

Next, we will cover the steps on how to create a log sink:

1. In the Cloud console, go to the Log Router, and select the Cloud project where you want to create the log sink.
2. Next, from the top left, select the **CREATE SINK** button and you will be presented with options as per *Figure 12.9*.
3. Here you can specify the details of the log sink, such as the sink name and its description; the name used in the example is `my-log-sink`.
4. Next, you can specify the destination where you want logs to be stored. There are different options available here, and based on your requirements, you can either select Google Cloud Storage (as per the example), create a new dataset in BigQuery or an existing dataset, specify Pub/Sub, or use **Splunk integration** as an option. In the example, we are using a Cloud Logging bucket, and the name of the bucket is `projects/Project_name/locations/australia-southeast1/buckets/my-regional-bucket`. You can either create a new bucket or use one that's already been created, such as the one we created in the previous step with an Australian region.

If your sink destination is a BigQuery dataset, the sink destination will be the following:
`bigrquery.googleapis.com/projects/PROJECT_ID/datasets/DATASET_`

5. Next, you can specify the inclusion and exclusion filter; this is again based on what you want to include or exclude from your logs. For simplicity, in this example, no inclusion or exclusion filters have been used.

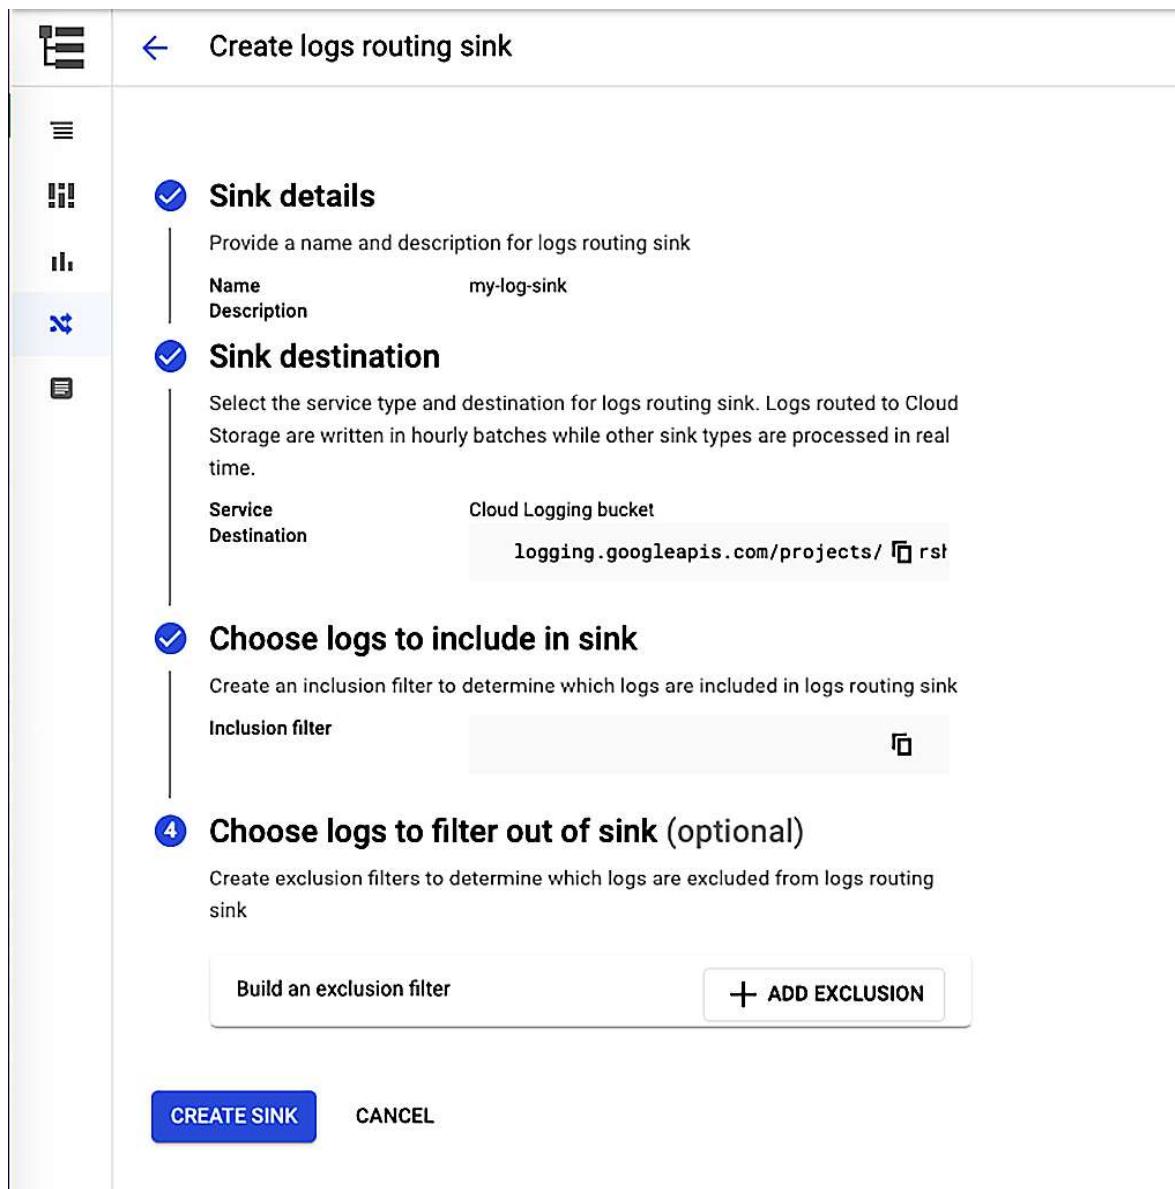


Figure 12.9 – Create a log sink

6. In *Figure 12.9*, we created a log sink named `my-log-sink` and selected the sink destination, its associated project, and geographical location, for example, `starbase155` and `australia-southeast1`, respectively. If you are routing logs between Google Cloud projects, you will need the appropriate destination permissions.
7. Finally, when all options are completed, click **CREATE SINK**.

Here we learned how to create an export and a log sink to send logs to the destination of your choice. Next, we will look at archiving and aggregating logs.

Log archiving and aggregation

The process of archiving logs is like the process outlined in the previous section. You create a log sink, specify the inclusion and exclusion criteria, and choose the Google Cloud Storage bucket as the destination.



Figure 12.10 – Log archiving pipeline

Figure 12.10 shows an illustration of what the log archiving pipeline looks like. You collect the events using Cloud Logging and then create a log sink with Cloud Storage as the destination.

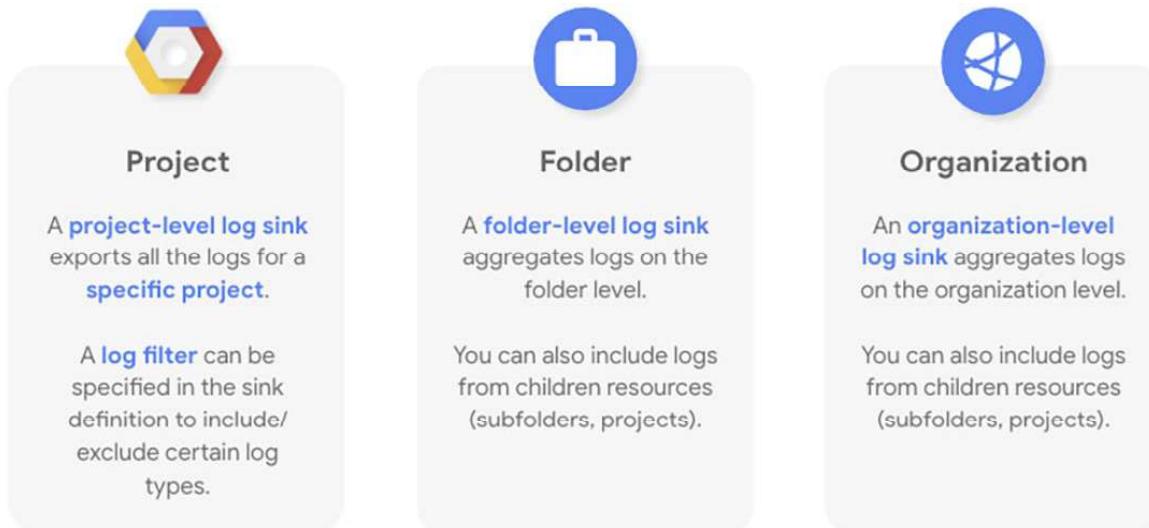


Figure 12.11 – Log aggregation

Log centralization is a typical requirement for auditing, archiving, and ensuring the non-repudiation of data and is considered a best practice. In order to achieve it, it is possible to configure aggregate exports at the organization level. This ensures that all logs on all projects (new and existing) in the organization are exported to a central location. *Figure 12.11* shows how you can configure log aggregation based on your requirements. You can aggregate logs at the project, folder, or organization level.